

# Design and Implementation of Vectorized Pseudorandom Number Generators

Master's Thesis Defense and Presentation

Markus Pawellek

May 21, 2020

# Outline

Introduction

Pseudorandom Number Generators

Design of the Library

Vectorization and SIMD Architectures

Implementation

- Xoroshiro128+

- MT19937

- Uniform Distribution Functions

Evaluation and Results

Conclusions and Future Work

# Introduction

# Introduction

What do we need random numbers for?

# Introduction

What do we need random numbers for?

- ▶ Physical Simulations, based on Monte-Carlo Methods

# Introduction

What do we need random numbers for?

- ▶ Physical Simulations, based on Monte-Carlo Methods



# Pseudorandom Number Generators

# True Randomness

What is a random sequence?



# True Randomness

What is a random sequence?

- ▶ existing formal concepts not applicable to computer systems

# True Randomness

What is a random sequence?

- ▶ existing formal concepts not applicable to computer systems
- ▶ nondeterministic, noncomputable, unpredictable

# True Randomness

What is a random sequence?

- ▶ existing formal concepts not applicable to computer systems
- ▶ nondeterministic, noncomputable, unpredictable
- ▶ generated by hardware components based on chaotic processes

# True Randomness

What is a random sequence?

- ▶ existing formal concepts not applicable to computer systems
- ▶ nondeterministic, noncomputable, unpredictable
- ▶ generated by hardware components based on chaotic processes

Disadvantages:

# True Randomness

What is a random sequence?

- ▶ existing formal concepts not applicable to computer systems
- ▶ nondeterministic, noncomputable, unpredictable
- ▶ generated by hardware components based on chaotic processes

Disadvantages:

- ▶ Unreproducibility

# True Randomness

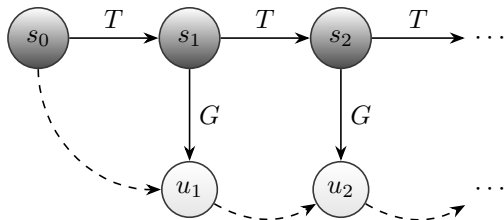
What is a random sequence?

- ▶ existing formal concepts not applicable to computer systems
- ▶ nondeterministic, noncomputable, unpredictable
- ▶ generated by hardware components based on chaotic processes

Disadvantages:

- ▶ Unreproducibility
- ▶ Speed Limitations

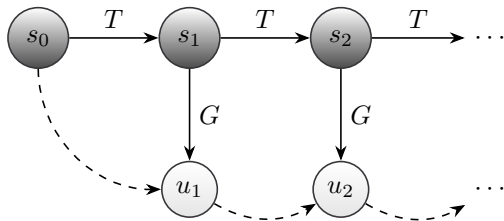
# Pseudorandom Number Generator Definition



$$\mathcal{G} := (S, T, U, G), \quad T: S \rightarrow S, \quad G: S \rightarrow U$$

$$s_0 \in S, \quad s_{n+1} := T(s_n), \quad u_n := G(s_n)$$

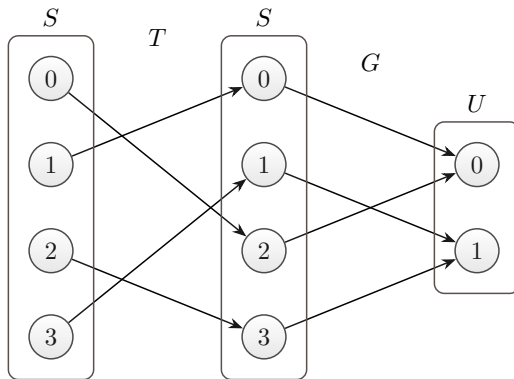
# Pseudorandom Number Generator Concept



$$s_0 \sim \mathcal{U}_S, \quad u_1 \leftarrow \mathcal{G}(), \quad u_2 \leftarrow \mathcal{G}(), \quad u_3 \leftarrow \mathcal{G}(), \quad \dots$$



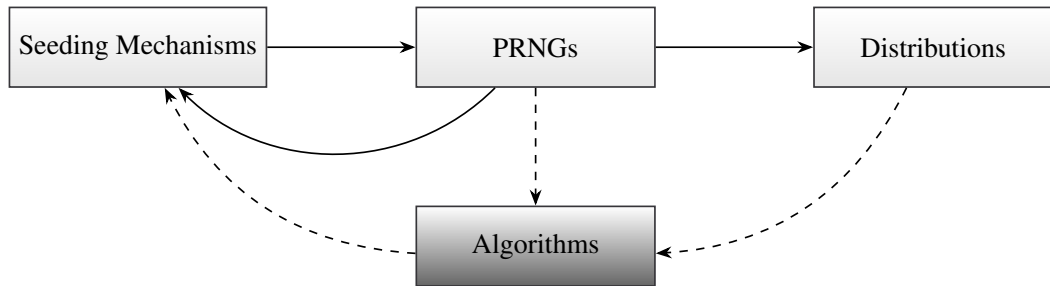
# Pseudorandom Number Generator Example



$$s_0 := 0, \quad (s_n) = \overline{2310}, \quad (u_n) = \overline{0110}$$

# Design of the Library

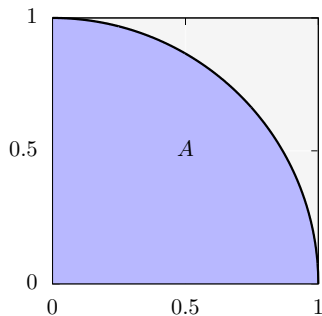
# Design Components



# Usage in C++

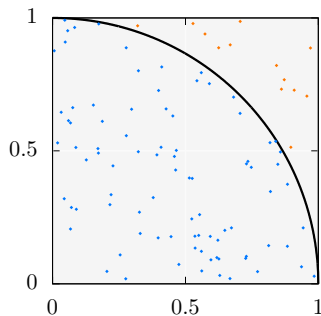
```
#include <pxart/pxart.hpp>
//
std::random_device rd{};
//
pxart::mt19937 rng1{};
pxart::mt19937 rng1{rd};
pxart::mt19937 rng1{pxart::mt19937::default_seeder{rd()}};
//
pxart::xrsr128p rng2{rng1};
//
const auto x = pxart::uniform<float>(rng1);
//
const auto y = pxart::uniform(rng2, -1.0f, 1.0f);
```

# Computation of $\pi$



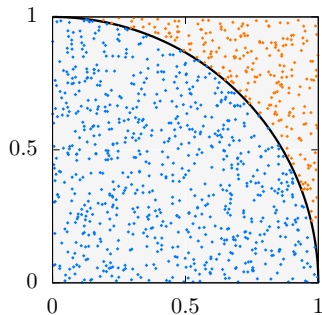
$$A = \frac{\pi}{4}, \quad \hat{\pi} = \frac{4N_A}{N}$$

# Computation of $\pi$



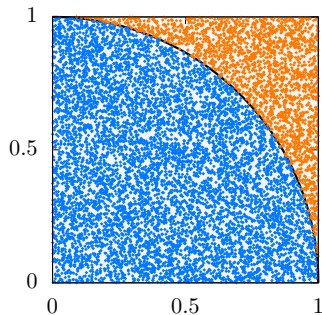
$$A = \frac{\pi}{4}, \quad \hat{\pi} = \frac{4N_A}{N} = \frac{4 \cdot 87}{100} = 3.48$$

# Computation of $\pi$



$$A = \frac{\pi}{4}, \quad \hat{\pi} = \frac{4N_A}{N} = \frac{4 \cdot 765}{1000} = 3.06$$

# Computation of $\pi$



$$A = \frac{\pi}{4}, \quad \hat{\pi} = \frac{4N_A}{N} = \frac{4 \cdot 7856}{10000} = 3.1424$$

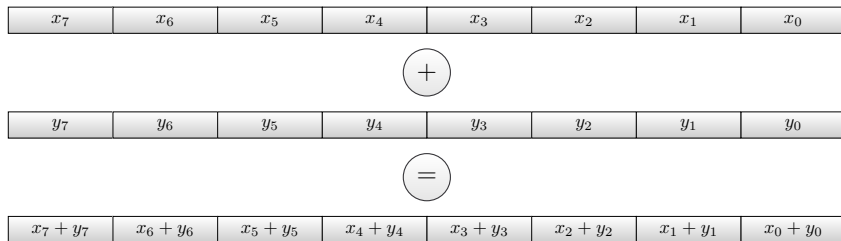


# Example Usage

```
// ...  
#include <pxart/pxart.hpp>  
// ...  
pxart::mt19937 rng{};  
const int samples = 100000000;  
int pi = 0;  
for (auto i = samples; i > 0; --i) {  
    const auto x = pxart::uniform<float>(rng);  
    const auto y = pxart::uniform<float>(rng);  
    pi += (x * x + y * y <= 1);  
}  
pi = 4.0f * pi / samples;  
  
// ...
```

# Vectorization and SIMD Architectures

# SIMD Architecture



- ▶ exploits data-level parallelism
- ▶ Intel CPUs use fixed-length vector registers
- ▶ vector operations are performed independently on all contained values at once

# SIMD Features in C++

- ▶ SSE (128-bit registers) and AVX (256-bit registers) instruction set features
- ▶ Assembler vs. Automatic Vectorization vs. Manual Vectorization by Intrinsics

```
// 128-bit registers
```

```
__m128 a;  
__m128d b;  
__m128i c;
```

```
c = _mm_add_ps(a, b);
```

```
// 256-bit registers
```

```
__m256 a;  
__m256i b;  
__m256d c;
```

```
c = _mm256_add_ps(a, b);
```

# SIMD Architecture

Why should we vectorize PRNGs manually?

- ▶ exploit full functionality of today's processors
- ▶ no automatic vectorization possible
- ▶ other vectorized code needs vectorized random numbers
- ▶ faster generation of numbers
- ▶ PRNGs are low-level, SIMD is low-level


# SIMD Architecture

What are conditions for good vectorization?

- ▶ no data dependency
- ▶ same processing pipeline
- ▶ branchless execution

# SIMD Example

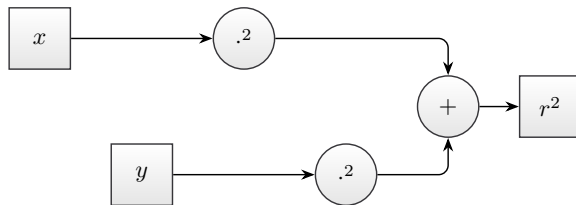
$$x, y \in \mathbb{R}, \quad r^2 = x^2 + y^2$$



```
double x = pxart::uniform<double>(rng);  
double y = pxart::uniform<double>(rng);
```

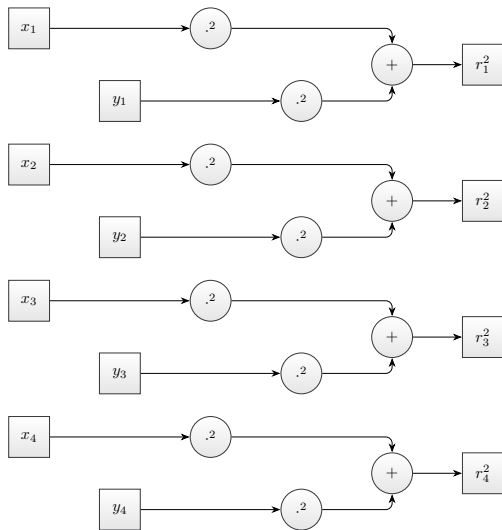
```
double x2 = x * x;  
double y2 = y * y;  
double r2 = x2 + y2;
```

# SIMD Example

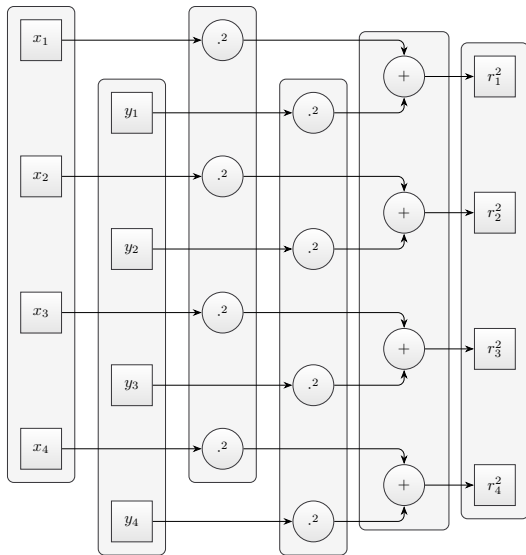




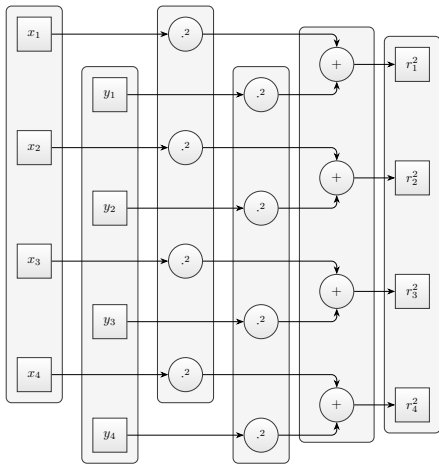
# SIMD Example



# SIMD Example




# SIMD Example



```
__m256d x = pxart::uniform<double>(vrng);  
__m256d y = pxart::uniform<double>(vrng);  
  
__m256d x2 = _mm256_mul_pd(x, x);  
__m256d y2 = _mm256_mul_pd(y, y);  
__m256d r2 = _mm256_add_pd(x2, y2);
```

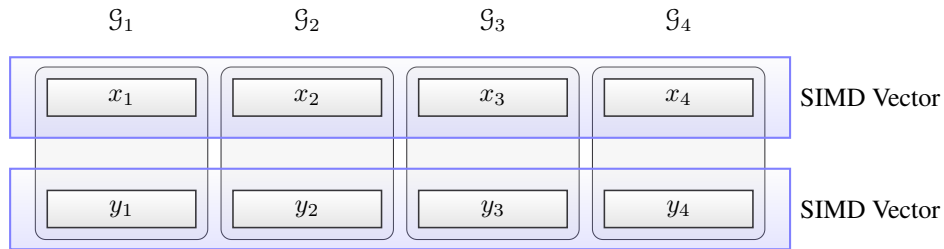
# Implementation

# Xoroshiro128+



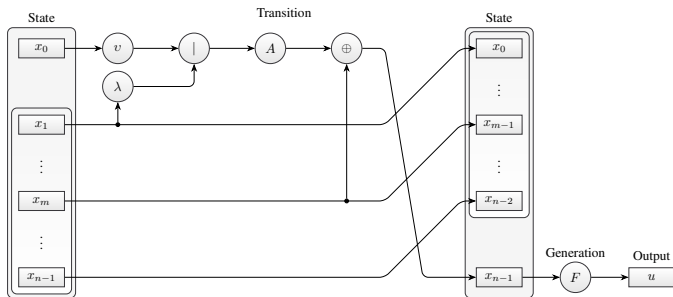

- ▶ scrambled linear PRNG
- ▶ 128-bit state, 64-bit output
- ▶ period:  $2^{128} - 1$
- ▶ jump operations

# Xoroshiro128+ Scalar and Vectorized




- ▶ several parallelization techniques
- ▶ multiple instances of the same generator
- ▶ seeding variations

# Mersenne Twister MT19937

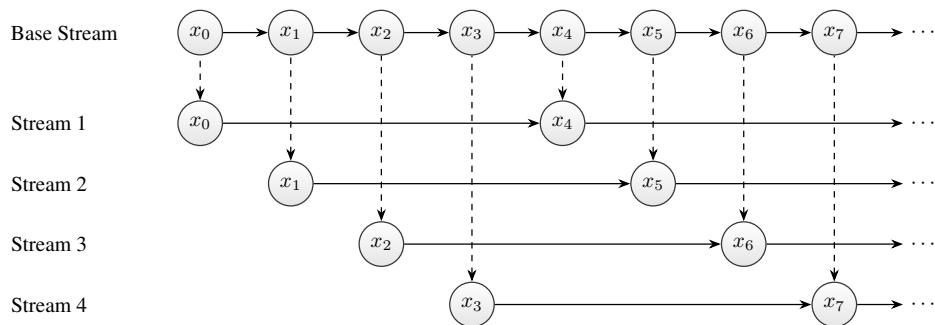


- ▶ de-facto standard
- ▶ linear PRNG
- ▶ 19937-bit / 19968-bit  $\approx$  2.4 KiB state, 32-bit output
- ▶ period:  $2^{19937} - 1$
- ▶ 623-dimensional equidistributed

# MT19937 Scalar








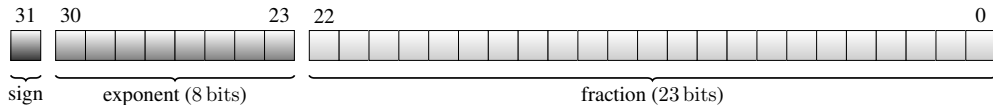
- ▶ vectorized generator will give same output as scalar one, only faster

# MT19937 SIMD



- implementation could be tested for the same output as the standard

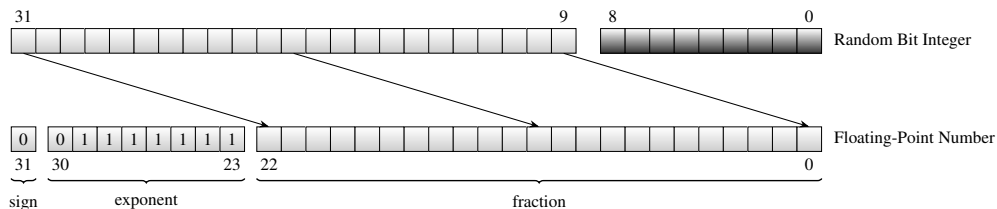
# Real Uniform Distribution: Floating-Point Encoding



$$x = (-1)^s \cdot m \cdot 2^{e-o}$$

- ▶ IEEE 754
- ▶ we use only normalized numbers

# Real Uniform Distribution



- ▶ get random integer
- ▶ shift bits with highest entropy into fraction part
- ▶ set sign and exponent to generate random floating-point value in  $[1, 2)$
- ▶ subtract one
- ▶ due independent operations easily vectorizable


# Integer Uniform Distribution

$$x \in \mathbb{N}_0, x < 2^{32}, \quad y = \left\lfloor \frac{(b-a) \cdot x}{2^{32}} \right\rfloor + a$$



- ▶ only approximation possible with vectorization
- ▶ first use bound by next more accurate multiplication
- ▶ then add offset

## Evaluation and Results

# Evaluation and Results



- ▶ Tests: TestU01, dieharder, Unit Tests, API Tests
- ▶ Performance: Benchmarks





## Conclusions and Future Work

# Comparison

	pXart	RNGAVXLIB	Intel MKL
Portable	✓	✗	✗
User-Friendly API	✓	✗	✗
Header-Only	✓	✗	✗
Open Source	✓	✓	✗
Documentation	✓	✗	✓
Distributions	✗	✓	✓
CMake and build2 Support	✓	✗	✗
Dependency-Free	✓	?	?


# Conclusions and Future Work

- ▶ possible applications in simulations
- ▶ mt19937 vs. xoroshiro128+

Thank you for Your Attention!

# References

# Processor



# Memory Hierarchy

