

Friedrich-Schiller-Universität Jena
Physikalisch-Astronomische Fakultät

**Design and Implementation of
Vectorized Pseudorandom Number Generators
and their Application to Simulation in Physics**

MASTER'S THESIS

for obtaining the academic degree

Master of Science (M.Sc.) in Physics

submitted by Markus Pawellek

born on May 7th, 1995 in Meiningen
Student Number: 144645

Primary Reviewer: Bernd Brüggemann

Primary Supervisor: Joachim Gießen

Jena, August 16, 2019

Abstract

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum. Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Acknowledgements

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Contents

Contents	i
List of Figures	iii
List of Abbreviations	v
Symbol Table	vii
1 Introduction	1
2 Background	3
2.1 Mathematical Preliminaries	3
2.2 Pseudorandom Number Generators	6
2.3 SIMD-Capable Processors	7
2.4 Simulation in Physics and Mathematics	7
2.5 Summary	7
3 Previous Work	9
3.1 The C++ API and Further Progressions	9
3.2 Techniques for Vectorization and Parallelization	9
3.3 Summary	9
4 Design of the API	11
5 Testing Framework	13
6 Implementation of Vectorized PRNGs	15
6.1 Linear Congruential Generators	15
6.2 Mersenne Twister	15
6.3 Permuted Congruential Generators	15
6.4 Xoroshiro	15
6.5 Middle Square Weyl Generator	15
6.6 Summary	15
7 Application to Simulations	17
8 Evaluation and Results	19
9 Conclusions	21
References	23

List of Figures

List of Abbreviations

Abbreviation	Definition
RNG	Random Number Generator
PRNG	Pseudorandom Number Generator
LCG	Linear Congruential Generator
MT	Mersenne Twister
MT19937	Mersenne Twister with period $2^{19937} - 1$
PCG	Permuted Linear Congruential Generator
CPU	Central Processing Unit
GPU	Graphics Processing Unit
SIMD	Single Instruction, Multiple Data
SSE	Streaming SIMD Extensions
AVX	Advanced Vector Extensions

Symbol Table

Symbol	Definition
$x \in A$	x ist ein Element der Menge A .
$A \subset B$	A ist eine Teilmenge von B .
$A \cap B$	$\{x \mid x \in A \text{ und } x \in B\}$ für Mengen A, B — Mengenschnitt
$A \cup B$	$\{x \mid x \in A \text{ oder } x \in B\}$ für Mengen A, B — Mengenvereinigung
$A \setminus B$	$\{x \in A \mid x \notin B\}$ für Mengen A, B — Differenzmenge
$A \times B$	$\{(x, y) \mid x \in A, y \in B\}$ für Mengen A und B — kartesisches Produkt
\emptyset	$\{\}$ — leere Menge
\mathbb{N}	Menge der natürlichen Zahlen
\mathbb{N}_0	$\mathbb{N} \cup \{0\}$
\mathbb{R}	Menge der reellen Zahlen
\mathbb{R}^n	Menge der n -dimensionalen Vektoren
$\mathbb{R}^{n \times n}$	Menge der $n \times n$ -Matrizen
$f: X \rightarrow Y$	f ist eine Funktion mit Definitionsbereich X und Wertebereich Y
$\partial\Omega$	Rand einer Teilmenge $\Omega \subset \mathbb{R}^n$
σ	Oberflächenmaß
λ	Lebesgue-Maß
$\int_{\Omega} f \, d\lambda$	Lebesgue-Integral von f über der Menge Ω
$\int_{\partial\Omega} f \, d\sigma$	Oberflächen-Integral von f über der Menge $\partial\Omega$
∂_i	Partielle Ableitung nach der i . Koordinate
∂_t	Partielle Ableitung nach der Zeitkoordinate
∂_i^2	Zweite partielle Ableitung nach i
∇	$\begin{pmatrix} \partial_1 & \partial_2 \end{pmatrix}^T$ — Nabla-Operator
Δ	$\partial_1^2 + \partial_2^2$ — Laplace-Operator
$C^k(\Omega)$	Menge der k -mal stetig differenzierbaren Funktion auf Ω
$L^2(\Omega)$	Menge der quadrat-integrierbaren Funktionen auf Ω
$H^1(\Omega)$	Sobolevraum
$f _{\partial\Omega}$	Einschränkung der Funktion f auf $\partial\Omega$
$\langle x, y \rangle$	Euklidisches Skalarprodukt
$[a, b]$	$\{x \in \mathbb{R} \mid a \leq x \leq b\}$
(a, b)	$\{x \in \mathbb{R} \mid a < x < b\}$
$[a, b)$	$\{x \in \mathbb{R} \mid a \leq x < b\}$
$u(\cdot, t)$	Funktion \tilde{u} mit $\tilde{u}(x) = u(x, t)$
A^T	Transponierte der Matrix A
id	Identitätsabbildung
$a := b$	a wird durch b definiert
$f \circ g$	Komposition der Funktionen f und g
$\begin{vmatrix} a & b \\ c & d \end{vmatrix}$	Determinante der angegebenen Matrix
$\text{span}\{\dots\}$	Lineare Hülle der angegebenen Menge
$ A $	Anzahl der Elemente in der Menge A

1 Introduction

For various mathematical and physical problems, there exists no feasible, deterministic algorithm to solve them. Especially, the simulation of physical systems with many coupled degrees of freedom, such as fluids, seem to be difficult to compute due to their high dimensionality. Instead, a class of randomized algorithms, called Monte Carlo methods, are used to approximate the actual outcome. Monte Carlo methods rely on repeated random sampling to obtain a numerical result. Hence, they are not bound to the curse of dimensionality and are able to evaluate complex equations quickly.

To obtain precise answers with a small relative error, Monte Carlo algorithms have to use a tremendous amount of random numbers. But the usage of truly random numbers generated by physical processes consists at least of two drawbacks. First, the output of the algorithm will be non-deterministic and, as a result, untestable. Second, the generation of truly random numbers is typically based on a slow process and consequently reduces the performance of the entire program. For that reason, Monte Carlo algorithms usually use so-called pseudorandom number generators. PRNGs generate a sequence of numbers based on a deterministic procedure and a truly random initial value as seed. The sequence of numbers is not truly random but fulfills several properties of truly random sequences.

The structure of Monte Carlo methods causes a program to spend most of its time with the construction of random numbers. Even the application of PRNGs does not change that. Today's computer processors provide functionality for the parallel execution of code in different ways, mainly SIMD and MIMD. Hence, to efficiently use the computing power of a CPU for Monte Carlo algorithms PRNGs have to be vectorized and parallelized to exploit such features. Whereas parallelization takes place at a high level, vectorization has to be done by the compiler or manually by the programmer at a much lower level. The implementation of PRNGs constraints automatic vectorization due to internal flow and data dependencies. To lift this restriction, a manual vectorization concerning data dependence and latencies appears to be the right way.

The C++ programming language is a perfect candidate for the development of vectorized PRNGs. It is one of the most used languages in the world and can be applied to small research projects as well as large enterprise programs. The language allows for the high-level abstraction of algorithms and structures. On the other hand, it is capable of accessing low-level routines to exploit special hardware features, like SSE, AVX, and threads. A typical C++ compiler is able to optimize the code with respect to such features automatically. But we as programmers are not bound to this and can manually optimize the code further. Every three years, a new standard is published, such as the new C++20 language specification. The language is evolving by its communities improvements and therefore it keeps to be a modern language. On top of this, other languages, such as Python, usually provide an interface to communicate with the C programming language. Through the design of an efficient implementation in C++, we can easily add support for other languages as well by providing a standard C interface.

Lots of PRNGs have been implemented by different libraries with different APIs. For

example, STL, Boost, Intel MKL, RNGAVXLIB, Lemire, tinyrng,... STL, Boost and ... provide a large set of robust PRNGs which are not vectorized but well documented. Their API makes them likely to be used but shows many flaws. It does not allow to explicitly use the vectorization capabilities of a PRNG, gives you a bad default seeding and makes use of standard distributions difficult and not adjustable. Lemire and RNGAVXLIB provide open-source, vectorized implementations with bad documentation and difficult-to-use code. Intel MKL as well provides vectorized PRNGs but is not available open-source and uses difficult interfaces. There is not any easily-accessible, portable, open-source library which gives a coherent, easy-to-use and consistent interface for vectorized PRNGs.

In this thesis, we develop a new library, called pxart, in the C++ programming language. pxart vectorizes a handful of already known PRNGs which partly do not exist as vectorized versions and provides a new API for their usage to accommodate the disadvantages of the standard random library of the STL. The library itself is header-only, open-source, and can be found on GitHub. It is easily installable on every operating system. Additionally, we compare the performance of our vectorized PRNGs to other already accessible implementations in Boost, Intel MKL, Lemire, RNGAVXLIB and others. The performance is measured by speed, code size, memory size, complexity, and random properties. Meanwhile, we apply the implementations to an example Monte Carlo simulation. For this, a small test framework is implemented which allows us to easily test and evaluate PRNGs with respect to stated measures.

2 Background

To systematically approach the implementation of PRNGs, basic knowledge in the topics of stochastics and finite fields is administrable. Together, these topics will give a deeper understanding of randomness in deterministic computer systems, a formal description of pseudorandom sequences and generators, and the mathematical foundation of Monte Carlo algorithms. Based on them, we are capable of scientifically analyzing PRNGs concerning their randomness properties. Vectorization techniques can be conceptualized by the architecture of modern SIMD-capable multiprocessors and their instruction sets. Especially the knowledge of typical instructions will make the design of a new API and its application to Monte Carlo simulations clear. The following sections will give an overview of the named topics.

2.1 Mathematical Preliminaries

Probability Theory

The observation of random experiments resulted in the construction of probability theory. But probability theory itself does not use a further formalized concept of randomness (Schmidt 2009). In fact, it allows us to observe randomness without defining it (Volchan 2002). Hence, we will postpone an examination of truly random sequences to the next section.

According to Schmidt 2009, Kolmogorov embedded probability theory in the theory of measure and integration. Albeit it heavily relies on measure-theoretical structures, probability theory is one of the most important applications of measure and integration theory. Therefore we will assume basic knowledge in this topic and refer to Schmidt 2009 and Elstrodt 2011 for a more detailed introduction to measure spaces, measurable functions, and integrals. Propositions and theorems will be given without proof.

The underlying structure of probability theory, which connects it to measure theory, is the probability space. It is a measure space with a finite and normalized measure. This gives access to all the usual results of measure theory and furthermore unifies discrete and continuous distributions. (Schmidt 2009, p. 193 ff.)

DEFINITION 2.1: (Probability Space)

A probability space is a measure space (Ω, \mathcal{F}, P) such that $P(\Omega) = 1$. In this case, we call P the probability measure, \mathcal{F} the set of all events, and Ω the set of all possible outcomes of a random experiment.

Due to the complex definition of a measure space, it is convenient to not have to explicitly specify the probability space when analyzing random experiments. Instead, we use random variables which are essentially measurable functions on a probability space (Schmidt 2009, p. 194). For complicated cases, these will serve as observables for specific properties and will make the analysis much more intuitive.

DEFINITION 2.2: (Random Variable)

Let (Ω, \mathcal{F}, P) be a probability space and (Σ, \mathcal{A}) a measurable space. A measurable function $X: \Omega \rightarrow \Sigma$ is called a random variable.

In this case, we denote with $P_X := P \circ X^{-1}$ the distribution and with $(\Sigma, \mathcal{A}, P_X)$ the probability space of X . Two random variables are identically distributed if they have the same distribution. Additionally, we say that X is a real-valued random variable if $\Sigma = \mathbb{R}$ and $\mathcal{A} = \mathcal{B}(\mathbb{R})$.

From now on, if a random variable is defined then, if not stated otherwise, it is assumed there exists a proper probability space (Ω, \mathcal{F}, P) and measurable space (Σ, \mathcal{A}) .

Another important concept of stochastics is known as independence. In Schmidt 2009 it is defined for a family of events, a family of sets of events, and a family of random variables. If we think of random variables as observables then their independence means that their outcomes do not influence each other. For our purposes, the general definition of all three forms of independence is distracting. In a computer, it makes no sense to talk about infinite sequences. Therefore the following definition of independence takes only a finite sequence of random variables into account. Furthermore, to make it more understandable, this definition uses a theorem from Schmidt 2009, p. 238 which characterizes the independence of random variables.

DEFINITION 2.3: (Independence)

Let $n \in \mathbb{N}$ and X_i be a random variable for all $i \in \mathbb{N}$ with $i \leq n$. We denote the respective random vector with $X := (X_i)_{i=1}^n$. Then these random variables are independent if the following equation holds.

$$P_X = \bigotimes_{i=1}^n P_{X_i}$$

Typical observations of random sequences include the estimation of the expectation value and the variance. Both of these values are needed for analyzing PRNGs and the development of Monte Carlo simulations (Landau and Binder 2015, p. 30 ff.). Due to their deep connection to the integral, both of these moments are defined for real-valued random variables. We give the usual definitions based on Schmidt 2009, p. 274 ff. in a simplified form.

DEFINITION 2.4: (Expectation Value and Variance)

Let X be a real-valued random variable such that $\int_{\Omega} |X| \, dP < \infty$. Then the

expectation value $\mathbb{E}X$ and variance $\text{var } X$ of X is defined in the following way.

$$\mathbb{E}X := \int_{\Omega} X(\omega) \, dP(\omega) , \quad \text{var } X := \mathbb{E} (X - \mathbb{E}X)^2$$

To not rely on the underlying probability space directly, we want to be able to compute the expectation value through the respective distribution of the random variable. The theory of measure and integration gives the following proposition, also known as rule of substitution (Schmidt 2009, p. 276).

PROPOSITION 2.1: (Substitution)

Let X be real-valued random variable and $f: \mathbb{R} \rightarrow \mathbb{R}$ a measurable function such that $\int_{\Omega} |f| \, dP_X < \infty$. Then the following equation holds.

$$\mathbb{E}(f \circ X) = \int_{\mathbb{R}} f(x) \, dP_X(x)$$

In particular, if $\mathbb{E}|X| < \infty$ then the above equation can be reformulated as follows.

$$\mathbb{E}X = \int_{\mathbb{R}} x \, dP_X(x)$$

The distribution of real-valued random variables is univariate and as a result can be described by so-called cumulative distribution functions (CDFs). The CDF intuitively characterizes the distribution and simplifies the analysis. Further, it can be proven that every CDF belongs to a univariate distribution. According to Schmidt 2009, p. 246, this is the theorem of correspondence. Sometimes it is even possible to define a probability density; a function that is the Lebesgue density of the respective distribution (Schmidt 2009, p. 255).

DEFINITION 2.5: (Probability Density and Cumulative Distribution Function)

Let X be a real-valued random variable. Then the respective cumulative distribution function is defined as follows.

$$F_X: \mathbb{R} \rightarrow [0, 1] , \quad F_X(x) := P_X((-\infty, x])$$

We call the function $p: \mathbb{R} \rightarrow [0, \infty)$ a probability density of X if for all $A \in \mathcal{B}(\mathbb{R})$

$$P_X(A) = \int_A p(x) \, d\lambda(x) .$$

As well as CDFs, probability densities can greatly simplify computations which are based on absolute continuous random variables. The following proposition, obtained from Schmidt 2009, shows the simplified computation of an expectation value through a Lebesgue integral.

PROPOSITION 2.2: (Chaining)

Let X be a real-valued random variable with probability density p . If $f: \mathbb{R} \rightarrow \mathbb{R}$ is a measurable function such that $\mathbb{E}|f \circ X| < \infty$ then

$$\mathbb{E}|f \circ X| = \int_{\mathbb{R}} f(x)p(x) d\lambda(x) .$$

A last important theorem to name is the strong law of large numbers (SLLN). According to Graham and Talay [2013](#), p. 13, the principles of Monte Carlo methods are based on this theorem. Please note, there exist many more variations of this theorem. We will again use a simplified version from Graham and Talay [2013](#).

THEOREM 2.3: (Strong Law of Large Numbers)

Let $(X_n)_{n \in \mathbb{N}}$ be a sequence of iid. real-valued random variables with finite expectation value μ . Then the following equation holds P -almost everywhere.

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n X_i = \mu$$

Finite Fields**2.2 Pseudorandom Number Generators**

PRNGs were first introduced by Neumann. A further discussion about randomness will make clear why the design of PRNGs can be seen as art.

Random and Pseudorandom Sequences

Originating in gambling and physical processes, randomness is a difficult concept and drives many philosophical discussions. Typically humans have a bad idea of randomness. Randomness makes only sense when it is applied to a sequence of values. Because we want to generate random numbers we only need a formal mathematical structure to work with. However, a widely accepted unique formal concept has not been found. But as stated in Volchan [2002](#) the question if a sequence is random decides at infinity. A random sequence, in general, is not computable or compressible by an algorithm. Even the methods to test this kind of randomness cannot be computed. For the development of RNGs on a computer, we cannot use such concepts. A computer in our sense is only capable of using finite sequences and has to compute its randomness to check it. Therefore, again as stated in Volchan [2002](#), we will stick to 'if it looks random, it is random'. In Kneusel [2018](#) the concept of randomness was simplified.

Random and Pseudorandom Number Generators**DEFINITION 2.6:** (Pseudorandom Number Generator)

A tuple (S, s_0, T, U, G) is called a PRNG. S is a non-empty, finite set of states. $s_0 \in S$ is the initial state. $T: S \rightarrow S$ is the transition function. U is a non-empty, finite set of output symbols. $G: S \rightarrow U$ is the output function which generates an output symbol for every state.

DEFINITION 2.7: (Pseudorandom Sequence of PRNG)

$(s_n)_{n \in \mathbb{N}}$ is the respective sequence of states

$$s_{n+1} := T(s_n)$$

Pseudorandom sequence $(u_n)_{n \in \mathbb{N}}$

$$u_n = G(s_n)$$

$$\begin{array}{ccccccc} s_0 & \xrightarrow{T} & s_1 & \xrightarrow{T} & s_2 & \xrightarrow{T} & \dots \\ & & G \downarrow & & \downarrow & & \\ & & u_1 & & u_2 & & \dots \end{array}$$

2.3 SIMD-Capable Processors**Architecture of Modern Central Processing Units****SIMD Instruction Sets and Efficiency****SSE, AVX, AVX512****2.4 Simulation in Physics and Mathematics****Mathematical and Physical Preliminaries****Baseline Model Problems****2.5 Summary**

3 Previous Work

3.1 The C++ API and Further Progressions

3.2 Techniques for Vectorization and Parallelization

3.3 Summary

The topic of PRNGs consists of several smaller parts. From a mathematical point of view, one has to talk about their definition and construction as well as methods on how to test their randomness. There have been a lot of publications concerning these issues. Hence, I am not able to give you a detailed overview. Instead, I will focus on the most relevant PRNGs and test suites, as well as some modern examples.

The creation of new PRNGs is sometimes understood to be black magic and can be hard since basically, one has to build a deterministic algorithm with a nearly non-deterministic output. In Kneusel 2018 one can find numerous different families of PRNGs. The most well-known ones are Linear Congruential Generators, Mersenne Twisters and Xorshift with its Variants. Whereas LCGs tend to be fast but weak generators in O'Neill 2014, one can find a further developed promising family of algorithms, called PCGs. Widynski 2019 describes another RNG based on the so-called middle square Weyl sequence. All of these generators have certain advantages and disadvantages in different areas such as security, games, and simulations.

After building a PRNG, one has to check if the generated sequence of random numbers fulfills certain properties. In general, these properties will somehow measure the randomness of our RNG. Typically, there are a lot of tests bundled inside a test suite such as TestU01 and Dieharder.

4 Design of the API

What do we want from the interface of our RNG? It should make testing with given frameworks like TestU01, dieharder, ent and PractRand easy. Benchmarking should be possible as well. Therefore we need a good API and a good application interface. Most of the time we want to generate uniform distributed real or integer numbers. We need two helper functions. So we see that the concept of a distribution makes things complicated. We cannot specialize distributions for certain RNGs. We cannot use lambda expressions as distributions. Therefore we want to use only helper functions as distributions and not member functions. So we do not have to specify a specialization and instead use the given standard but we are able to do it. Therefore functors and old-distributions are distributions as well and hence we are compatible to the standard.

Additionally, we have to be more specific about the concept of a random number engine. The output of a random number engine of the current concept is magical unsigned integer which should be uniformly distributed in the interval $[min, max]$. But these magic numbers can result in certain problems if used the wrong way, see Melissa O'Neill Seeding Surprises. Therefore the general idea is to always use the helper functions as new distributions which define min and max explicitly and make sure you really get those values. This is also a good idea for the standard. And it is compatible with the current standard.

Now think of vector registers and multiprocessors. The random number engine should provide ways to fill a range with random numbers such that it can perform generation more efficiently. Think about the execution policies in C++17. They should be provided as well.

5 Testing Framework

6 Implementation of Vectorized PRNGs

6.1 Linear Congruential Generators

6.2 Mersenne Twister

6.3 Permuted Congruential Generators

6.4 Xoroshiro

6.5 Middle Square Weyl Generator

6.6 Summary

7 Application to Simulations

8 Evaluation and Results

godbolt google benchmark intel vtune amplifier testu01 dieharder

9 Conclusions

References

General

Elstrodt, Jürgen: *Maß- und Integrationstheorie*. Siebte korrigierte und aktualisierte Auflage. Springer, 2011. ISBN: 978-3-642-17904-4.

Graham, Carl and Denis Talay: *Stochastic Simulation and Monte Carlo Methods: Mathematical Foundations of Stochastic Simulation*. Springer, 2013. ISBN: 978-3-642-39362-4.

Kneusel, Ronald T.: *Random Numbers and Computers*. Springer, 2018. ISBN: 978-3-319-776696-5.

Landau, David P. and Kurt Binder: *A Guide to Monte Carlo Simulations in Statistical Physics*. Fourth Edition. Cambridge University Press – University of Cambridge, 2015. ISBN: 978-1-107-07402-6.

O'Neill, Melissa E.: *PCG: A Family of Simple Fast Space-Efficient Statistically Good Algorithms for Random Number Generation*. In: *ACM Transactions on Mathematical Software* (2014).

Schmidt, Klaus D.: *Maß und Wahrscheinlichkeit*. Springer, 2009. ISBN: 978-3-540-89729-3.

Volchan, Sérgio B.: *What is a Random Sequence?* In: *The American Mathematical Monthly* 109.1 (2002), pp. 46–63.

Widynski, Bernard: *Middle Square Weyl Sequence RNG*. In: *arXiv preprint arXiv:1704.00358v3* (2019).

Statutory Declaration

I declare that I have developed and written the enclosed Master's thesis completely by myself, and have not used sources or means without declaration in the text. Any thoughts from others or literal quotations are clearly marked. The Master's thesis was not used in the same or in a similar version to achieve an academic grading or is being published elsewhere.

On the part of the author, there are no objections to the provision of this Master's thesis for public use.

Jena, August 16, 2019

Markus Pawellek